Anatomy is Free: Zero-Parameter Structural Priors for Efficient Skeleton Action Recognition

Anonymous Author(s)

Abstract

Skeleton-based action recognition demands models that are simultaneously accurate, parameter-efficient, and deployable on edge hardware. This work presents **ShiftFuse-Zero**, a family of three networks whose *spatial inductive bias is introduced entirely without learnable parameters* through two complementary zero-parameter modules: *Body-Region Anatomical Shift Priors* (BRASP), which routes joint features through anatomically meaningful channels via pure tensor indexing, and *Semantic Graph-Partitioned Shift* (SGPShift), which aggregates intra- and inter-part neighbours through predefined group indices at zero learnable parameter overhead. These free structural priors reduce the burden on downstream learnable modules, enabling the compact *EfficientZero Block* to operate on an already anatomically-refined manifold. Three models cover the edge-to-server spectrum: **Nano** (94 K), **Small** (281 K), and **Large** (1.18 M). On NTU RGB+D 60 X-Sub, results without distillation are 86.2 % (Nano), 88.8 % (Small), and 92.5 % (Large); on NTU-120 X-Sub, SFZ-Large reaches **88.5%**, matching EfficientGCN-B4 (88.3%) at 4.3× fewer FLOPs. Knowledge distillation (Large teacher) further lifts Nano/Small to 88.5 %/90.4 %; Small+KD matches or marginally surpasses EfficientGCN-B0 (90.2%) within a two-backbone design at 7× fewer FLOPs. To our knowledge, SFZ-Nano dominates the accuracy–efficiency frontier among sub-100 K parameter skeleton models, runs in **11 ms** on desktop CPU and **50 ms** on mobile (Snapdragon 8 Gen 3), and fits in **370 KB**; with knowledge distillation, Nano matches AGCN (6.9 M parameters) at 74× higher parameter efficiency.

CCS Concepts

• **Computing methodologies → Activity recognition and understanding**.

Keywords

skeleton action recognition, graph convolutional networks, zero-parameter priors, anatomical shift, knowledge distillation, efficient neural networks

ACM Reference Format:
Anonymous Author(s). 2026. Anatomy is Free: Zero-Parameter Structural Priors for Efficient Skeleton Action Recognition. In *Proceedings of the 34th ACM International Conference on Multimedia.* ACM, New York, NY, USA, 9 pages.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
MM '26, Melbourne, Australia
© 2026 ACM.

1 Introduction

Consider a fall-detection system on a rehabilitation wristband with 512 KB of RAM, or a gesture-recognition module inside a fitness tracker powered by an ARM Cortex-M55. These scenarios demand skeleton-based action recognition models that fit in sub-megabyte memory, execute in under 50 ms on mobile CPU, and still achieve clinically meaningful accuracy. Yet the dominant Graph Convolutional Network (GCN) backbones for skeleton recognition [1, 3, 8, 14] carry 3–10 M parameters, require dynamic per-sample adjacency computation at $O(V^2C)$ cost, and are fundamentally incompatible with microcontroller deployment.

Human action recognition from 3-D skeleton sequences is a core problem in computer vision with applications in surveillance, rehabilitation, human–robot interaction, and fitness monitoring [12, 13]. GCNs [7] have become the dominant spatial backbone since ST-GCN [18], with successive works adding adaptive adjacency matrices [14], multi-scale cross-spacetime graphs [11], channel-wise topology refinement [1], information-theoretic representation learning [3], and hierarchical body-part decomposition [8]. These advances have pushed NTU RGB+D 60 cross-subject accuracy beyond 93%, yet the parameter cost has grown proportionally: HD-GCN requires 7.5 M parameters and InfoGCN 9.9 M.

The efficient-architecture line [16, 19] tackles this disparity via compound scaling and depthwise convolutions, but largely ignores the rich *anatomical structure* of the human body as a source of *free* inductive bias. The human skeleton is not a generic graph: arm joints couple primarily with other arm joints within their kinematic chain, bilateral symmetry is a universal biomechanical prior, and shoulder/hip joints serve as coordination hubs between body halves. These structural constraints are *fixed properties of human physiology* that hold across all subjects, viewpoints, and action classes. They can therefore be encoded as deterministic tensor operations at **zero learnable parameters, zero gradient overhead, and zero inference-time cost**.

This paper presents the first systematic study of *zero-parameter* anatomically-guided inductive bias across a multi-scale skeleton GCN family — distinct from graph-distance shifts [2] (hop count, not body-part identity) and ST-GCN's fixed adjacency [18] (graph distance, not semantic group membership). HD-GCN and SkateFormer exploit anatomy through *learnable* grouping; no prior work has measured the accuracy contribution of purely fixed anatomical routing across Nano-to-Large scale. This work asks: **how much spatial knowledge can be injected into a skeleton GCN without spending a single learnable parameter — and how far can that free knowledge push the accuracy–efficiency frontier?**

Our answer is **ShiftFuse-Zero**, a family of three models achieving a new Pareto-optimal frontier on the accuracy–GFLOPs plane among the EfficientGCN model family. The core contribution is BRASP and SGPShift, which contribute +1.8 pp *at zero FLOPs cost*; the additional 7× FLOPs advantage over EfficientGCN-B0 is an

independent architectural effect arising from the two-backbone design and reduced per-backbone channel width (see Table 4 and Section 4.5). SFZ-Nano (94 K) and SFZ-Small (281 K) redefine edge-deployment with ultra-low FLOPs and quantisation-friendly operations. On NTU-60 X-Sub, SFZ-Nano reaches **86.2%** without any distillation and **88.5%** with knowledge distillation, matching AGCN's accuracy (6.9 M params) within 370 KB of storage (KD is an optional plugin requiring SFZ-Large as teacher; all primary results are without KD). SFZ-Small achieves **88.8%** at comparable parameter count (281 K vs. 290 K) with 7× fewer FLOPs than EfficientGCN-B0; with KD, Small+KD reaches **90.4%**, matching or marginally surpassing B0 while retaining the 7× FLOPs advantage. SFZ-Large pushes the absolute accuracy frontier to **92.5%** at 1.18 M parameters. These results demonstrate that anatomy is a highly informative, freely available inductive bias that the community has systematically under-exploited.

Contributions.

(1) **BRASP** — Body-Region Anatomical Shift Priors: routes joint features through four combined bilateral channel groups via a precomputed gather permutation; zero parameters, zero gradient (Section 3.3).
(2) **SGPShift** — Semantic Graph-Partitioned Shift: routes each channel's joint features to a ranked intra-part or inter-part neighbour via precomputed index buffers (identity channels unchanged); zero parameters (Section 3.3).
(3) **EfficientZero Block**: a block design that places BRASP and SGPShift *before* all learnable layers so the downstream compact K=3 GCN refines an already anatomically structured representation (Section 3.4).
(4) **ShiftFuse-Zero model family**: three architectures — Nano (94 K), Small (281 K), Large (1.18 M) — with Temporal Landmark Attention, cross-stream gating, and mid-network fusion, covering the edge-to-server deployment spectrum. Vanilla Hinton KD [6] serves as an optional plugin (requiring SFZ-Large as teacher), yielding up to +2.3 pp with no additional inference overhead (Section 3.5).
(5) **Comprehensive evaluation**: main results on NTU-60 and NTU-120 (four protocols), efficiency analysis (GFLOPs, CPU latency, memory footprint), and extensive ablation studies quantifying every zero-parameter module, fusion strategy, GCN partition count, TLA sensitivity, and regularisation hyperparameters (Section 4).

2 Related Work

GCN-based skeleton recognition. ST-GCN [18] established the skeleton graph and spatial-temporal convolution paradigm using a three-subset partitioning (self, centripetal, centrifugal) of the physical adjacency matrix. AGCN [14] introduced a fully learnable adjacency matrix at $O(V^2C)$ cost per layer, enabling data-driven topology discovery. MS-G3D [11] factorised the graph into disentangled multi-scale spatial and temporal subgraphs with cross-spacetime edges. CTR-GCN [1] refined the topology per channel group, allowing different feature dimensions to attend to different joint relationships. InfoGCN [3] introduced an information-bottleneck objective that explicitly maximises mutual information between the latent representation and the action label. DSTA-Net [15] decouples spatial and temporal self-attention. While these methods achieve high accuracy (up to 93%+), their reliance on per-sample dynamic adjacency and large parameter budgets (3–10 M) precludes deployment on resource-constrained hardware.

Body-part decomposition. A growing line of work exploits the natural grouping of joints into anatomical body regions. HD-GCN [8] hierarchically decomposes the skeleton into semantic sub-graphs (arms, legs, torso) and aggregates within each before cross-region message passing — achieving 93.4% on NTU-60 X-Sub but at 7.5 M parameters. BlockGCN [20] uses persistent homology from topological data analysis to encode bone connectivity at the block level. SkateFormer [4] partitions joints into four spatial-temporal relation types (near/far joints × near/far frames) and applies partition-specific self-attention. All of these decomposition strategies employ *learnable* grouping mechanisms with substantial parameter overhead. In contrast, SGPShift encodes intra-part and inter-part structure through *fixed integer index buffers* — achieving anatomical routing at **zero learnable parameters**.

Efficient skeleton networks. EfficientGCN [16] introduced compound scaling (α=1.2, β=1.35) and multiple input branches (MIB) to build a family of models from B0 (0.29 M) to B4 (1.1 M), establishing the efficiency frontier for skeleton GCNs. SGN [19] exploits semantic joint embeddings and frame-level dynamics gating at 0.69 M parameters. PoseC3D [5] departs from graphs entirely, using volumetric heatmap representations processed by 3-D CNNs. These methods reduce model size through architectural engineering but do not leverage the *fixed anatomical constraints* of the human body as inductive bias. ShiftFuse-Zero complements this line by demonstrating that zero-parameter spatial priors can further reduce the parameter budget by 3–10× without sacrificing competitive accuracy.

Temporal modelling in skeleton networks. Temporal reasoning has evolved from early LSTM-based encoders to the dedicated temporal convolution modules introduced by ST-GCN [18] and its successors. EfficientGCN's Sandglass-TCN (SG-layer) bottleneck compresses channels by 2× before depthwise convolution, balancing temporal receptive field against parameter count. MS-G3D introduced disentangled multi-scale temporal subgraphs. Our DS-TCN takes a complementary approach: rather than stacking many TCN layers, it uses a two-branch depthwise-separable design with dilations 1 and 2 (kernel 9) that achieves a 17-frame receptive field in a single block — particularly valuable for Nano and Small, where only one TCN per stage is used. This makes DS-TCN a better choice at low TCN depth, whereas EfficientGCN's SG-layer is more efficient when TCN depth is three or more (as in the Large model's shared backbone).

Shift operations and zero-parameter spatial mixing. The shift operation [17] was introduced in image CNNs to inject spatial structure by physically displacing channels along spatial dimensions — no multiply-accumulate, no learnable parameters. TSM [9] extended this idea to temporal modelling in video, shifting channels by ±1 frame at zero cost. Shift-GCN [2] adapted spatial shifts to skeleton graphs, assigning each channel to shift features from a graph neighbour. However, Shift-GCN assigns shift patterns by graph distance

(1-hop, 2-hop), ignoring the semantic meaning of body regions. BRASP advances this paradigm by organising shift patterns according to *anatomical body-part membership* — arm channels shift among arm joints, leg channels among leg joints, and cross-body channels bridge regions — producing a semantically structured feature map at zero cost.

Knowledge distillation. Hinton *et al.* [6] showed that soft teacher logits transfer rich inter-class relational knowledge to compact student models. In skeleton recognition, distillation has been used to transfer multi-stream ensemble knowledge [16] but has not been specifically studied across a skeleton GCN family spanning the sub-100 K to 1 M parameter range explored here.

3 Methodology

3.1 Overview

Figure 1 shows the end-to-end pipeline. A raw skeleton sequence is expanded into four complementary streams (joint, velocity, bone, bone-velocity); all four are pre-computed offline and each model selects its subset (Nano: three streams; Small/Large: all four). Selected streams are fused by a learned *StreamFusionConcat* stem, processed by a stack of *EfficientZero Blocks* (Nano/Small) or *B4Blocks* (Large), and a global pooling head (gated for Nano/Small, plain GAP for Large) produces class logits.

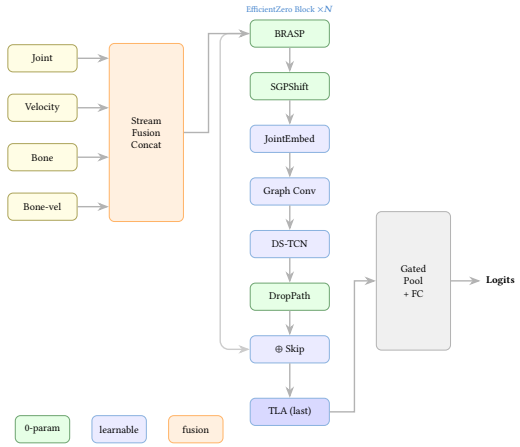

Figure 1: ShiftFuse-Zero end-to-end pipeline. Four streams are pre-computed; each model selects its subset (Nano: 3, Small/Large: 4) for fusion by a learnable stem. EfficientZero Blocks (green = zero-parameter) follow. A gated pooling head produces class logits.

3.2 Skeleton Streams

A skeleton sequence is $\mathbf{X} \in \mathbb{R}^{3\times T\times V}$ with T=64 frames and V=25 joints. This work constructs four complementary streams: joint positions $\mathbf{X}^{(j)}$, velocity $\mathbf{X}^{(v)}=\mathbf{p}_{t+1}-\mathbf{p}_t$, bone vectors $\mathbf{X}_v^{(b)}=\mathbf{p}^{(v)}-\mathbf{p}^{(\mathrm{par}(v))}$, and bone velocity $\mathbf{X}^{(bv)}=\Delta\mathbf{X}^{(b)}$ (where Δ denotes the first-order temporal difference). Frame-0 is padded by edge replication; downstream BN absorbs the resulting constant offset. *StreamFusionConcat* [16] applies a per-stream BN, concatenates (N_s×3 channels, $N_s \in \{3, 4\}$ depending on model), and projects to C_{stem} via a 1×1 convolution with learned per-stream weighting.

3.3 Zero-Parameter Spatial Modules

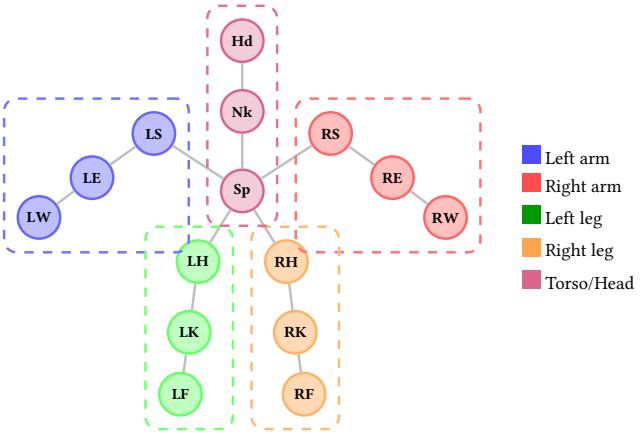

Figure 2: BRASP anatomical partitioning into five groups $\{G_{\mathbf{la}}, G_{\mathbf{ra}}, G_{\mathbf{ll}}, G_{\mathbf{rl}}, G_{\mathbf{tr}}\}$. Per-channel permutations σ_c (Eq. 1) are derived from these groups; no gradient flows through them.

BRASP.. Figure 2 illustrates the five physical joint regions. These five regions are combined into *four channel groups* by merging bilateral pairs: left-arm ∪ right-arm → *arm*; left-leg ∪ right-leg → *leg*; torso; cross-body. Given a feature map $\mathbf{F} \in \mathbb{R}^{C\times T\times V}$, BRASP assigns per-channel fixed gather permutations $\sigma_c : V \to V$:

$$\tilde{\mathbf{F}}_{c,t,v} = \mathbf{F}_{c,t,\sigma_c(v)}, \quad v \in \{0, \ldots, V-1\}, \tag{1}$$

Arm channels (25%) cycle within the combined {left-arm ∪ right-arm} joint set (12 joints), enabling direct bilateral feature exchange; *leg* (25%) within combined leg joints (8 joints); *torso* (12.5%); and *cross-body* (37.5%) cycles via graph-neighbour adjacency. Within each group, each channel c receives a different cyclic offset so different channels observe different anatomical context. All σ_c are precomputed integer buffers — **zero parameters, zero gradient**.

Remark 1 *(BRASP bridges graph-far symmetric joints).* Bilaterally symmetric pairs (*e.g.*, left-wrist / right-wrist) are separated by d_G=10+ hops in the NTU-25 skeleton and are invisible to the K=3 GCN in a single layer. BRASP's intra-part permutation directly places right-wrist features at left-wrist *before* any graph convolution, bridging the 10-hop graph distance at zero parameter cost. After one BRASP shift, joint v receives features from an anatomical neighbour $\sigma_c(v)$ that may be graph-far but biomechanically coupled, in a single index-lookup step.

SGPShift. Figure 3 illustrates the intra/inter channel routing. Like BRASP (Eq. 1), SGPShift is a per-channel gather permutation,

$$\hat{\mathbf{F}}_{c,t,v} = \mathbf{F}_{c,t,\pi_c(v)}, \tag{2}$$

where $\pi_c : V \to V$ routes each channel to a ranked neighbour from a precomputed graph partition. Channels split into three approximately equal groups ($C/3$ each, rounded to nearest integer): *intra-part* channels shift to the top-ranked same-part neighbour (via precomputed $\mathbf{A}_{\text{intra}}$); *inter-part* channels shift to the top-ranked cross-part boundary neighbour (via $\mathbf{A}_{\text{inter}}$); and *identity* channels pass through unchanged. All routing indices are precomputed integer buffers — unlike HD-GCN and SkateFormer [4, 8] which use

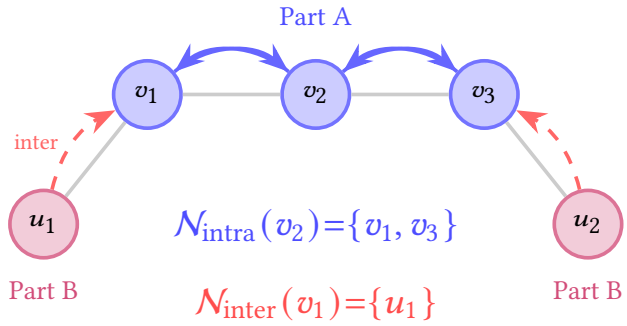

Figure 3: SGPShift aggregation. Blue solid arrows: intra-part neighbours. Red dashed arrows: inter-part boundary neighbours. All sets are precomputed integer buffers — zero learnable parameters.

learnable per-group parameters, SGPShift encodes this anatomical routing at **zero learnable parameters**.

3.4 EfficientZero Block

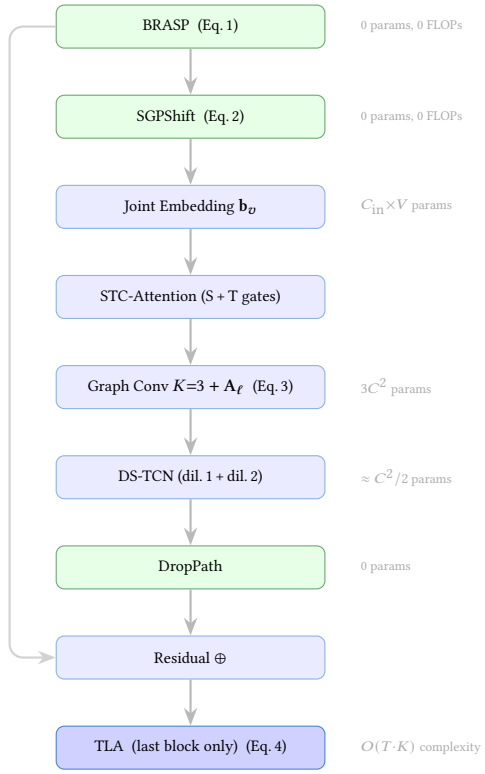

Figure 4: EfficientZero Block. Green = zero-parameter (BRASP, SGPShift, DropPath). Per-layer parameter counts shown on the right. TLA is only active at the last block of the final stage.

Figure 4 shows the full block pipeline with per-layer costs.

(1) **BRASP + SGPShift**: anatomical pre-structuring at zero cost.

(2) **Joint Embedding**: per-joint additive bias $\mathbf{b}\in\mathbb{R}^{C_{\text{in}}\times V}$ (SGN joint encoding [19]); no weight decay.

Table 1: ShiftFuse-Zero variant configurations. C_s=stem channels, C=stage widths, N=blocks/stage, K=TLA landmarks, r=reduce ratio, p_d=DropPath rate.

Model	C_s	C	N	K	r	p_d	FLOPs	Params
Nano	24	[32,64,128]	[1,1,1]	12	8	0.10	0.13 G	94 K
Small	24	[32,64,128]	[1,2,1]	8	8	0.10	0.39 G	281 K
Large†	64	see †	—	14	8	0.20	1.95 G	1.18 M

†B4-style: per-stream 64→96→48 (d 2+2), concat 144, shared 144→128→272 (d 3+3).

(3) **STC-Attention**: spatial ($\mathbf{A}_s\in\mathbb{R}^V$, mean 1) and temporal ($\mathbf{A}_t\in\mathbb{R}^T$, avg+max Conv_{1d} kernel 9) gates applied before the GCN via element-wise broadcast $\mathbf{F}'_{c,t,v}=\mathbf{F}_{c,t,v}\cdot[\mathbf{A}_s]_v\cdot[\mathbf{A}_t]_t$. Output: $(1-\sigma(\gamma))\mathbf{F}+\sigma(\gamma)\mathbf{F}'$, $\gamma_0=-4$.

(4) **Graph Convolution** (K=3): ST-GCN [18] partition adjacency plus a single shared learned matrix $\mathbf{A}_\ell$:

$$\mathbf{Z}=\sum_{k=0}^{2}\mathbf{W}_k\big(\mathbf{A}_k+|\mathbf{A}_\ell|\big)\mathbf{F}, \tag{3}$$

$|\cdot|$ enforces non-negative weights and allows $\mathbf{A}_\ell$ to create new connections (where $\mathbf{A}_\ell<0$ trained, $|\mathbf{A}_\ell|>0$ yields a positive edge), making it strictly more expressive than a pure non-negativity clamp; $|\cdot|$ also has no zero plateau. $\mathbf{A}_\ell\in\mathbb{R}^{V\times V}$ is shared (globally across all blocks in Nano; within-stage in Small), 625 params per shared set, no weight decay.

(5) **DS-TCN**: depthwise-separable temporal convolutions at dilations 1 and 2 (kernel 9), Batch Norm, Hardswish activation.

(6) **DropPath** (regularisation): stochastic depth; 0 parameters.

(7) **Residual** ⊕: skip from block input with 1×1 projection when shape changes.

Block Ordering Matters. Placing BRASP and SGPShift *before* the GCN is essential: the GCN refines an already anatomy-structured representation. Moving them *after* the GCN costs −1.9 pp (Table 8), confirming pre-structuring must precede message passing.

Complementary Temporal Module: TLA.. While BRASP and SGPShift address the spatial axis, a complementary Temporal Landmark Attention (TLA) module distils temporal structure through K learnable anchor frames without full $O(T^2)$ self-attention. TLA applies the deformable-attention principle [21] to the temporal axis: anchor positions are continuous and gradient-trainable via bilinear interpolation, reducing attention complexity from $O(T^2)$ to $O(T{\cdot}K)$:

$$\tau_k=T\cdot\sigma(\alpha_k),\quad k=1,\dots,K, \tag{4}$$
$$\mathbf{m}_t=\sum_{k=1}^{K}\text{Bilin}(\mathbf{F},\tau_k)\cdot\text{softmax}_k\big(\langle\mathbf{q}_t,\mathbf{k}_{\tau_k}\rangle\big),$$

where α_k are learned logits; $\text{Bilin}(\mathbf{F},\tau_k)$ interpolates $\mathbf{F}$ along the T-axis at fractional position $\tau_k\in[0,T-1]$, yielding a $\mathbb{R}^{C\times V}$ feature vector that enables gradient flow to α_k; and $\mathbf{q}_t,\mathbf{k}_{\tau_k}\in\mathbb{R}^{C_r}$ are projections, $C_r=C/r$. Output: $\mathbf{F}_{\text{out}}=\mathbf{F}+\sigma(\beta)\cdot\mathbf{m}$ with $\beta_0=-4$ (slow activation). Complexity $O(T{\cdot}K)$ vs. $O(T^2)$ for full attention.

3.5 Model Family

Table 1 lists the key hyperparameters per variant.

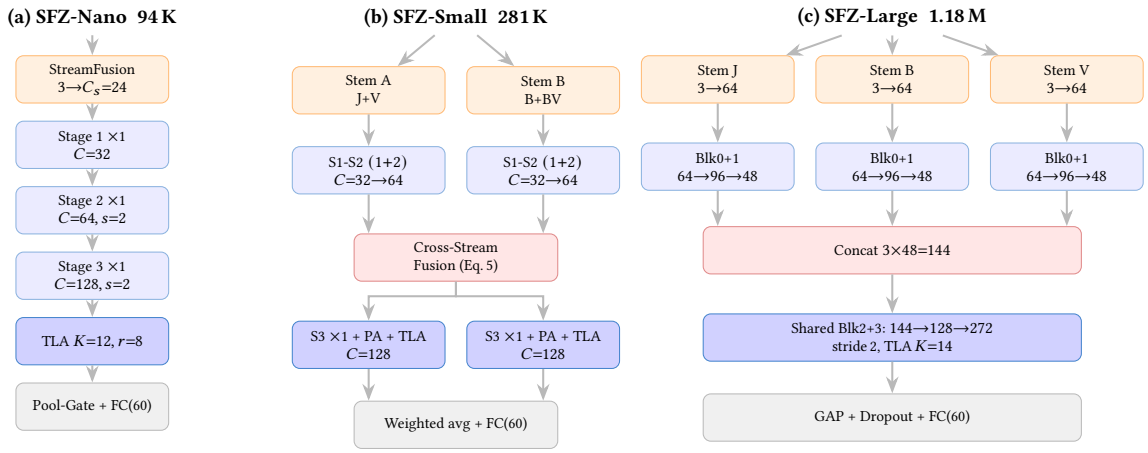

Figure 5: ShiftFuse-Zero model family. (a) Nano: single backbone, 3-stream early fusion, TLA (K=12). (b) Small: two-backbone late fusion with Cross-Stream Fusion, PartAttention (PA), TLA (K=8). (c) Large: B4-style mid-network fusion — per-stream Blk0+1 (64→96→48), concat 144 ch, shared Blk2+3 (144→128→272, TLA K=14).

Nano (94 K).. Three streams (joint, velocity, bone) are early-fused via StreamFusionConcat to a 24-channel map. A single backbone [32, 64, 128] with one block per stage processes it. TLA (K=12, r=8, C_r=16) adds only 8.2 K parameters; channel SE is omitted to keep the model within the 100 K budget.

Small (281 K).. Two independent backbones process two stream pairs: joint+velocity (Backbone A) and bone+bone-velocity (Backbone B). Stage 2 uses tcn_depth=2 for richer temporal modelling; the final stage adds *PartAttention* [16] (pa variant, in EfficientGCN's code but unused in B0/B4): per-group GAP over five body-part groups, bottleneck-projected ($C \to C/16 \to 1$) to per-joint scalars (12.3 K params). After stage 2, *Cross-Stream Fusion* exchanges information:

$$\mathbf{F}' = \mathbf{F} + \sigma(\beta) \cdot W_{\uparrow}\, \mathrm{ReLU}\big(W_{\downarrow}\,[\mathbf{F}_A; \mathbf{F}_B]\big), \tag{5}$$

$W_{\downarrow} : 2C \to C/4$, $W_{\uparrow} : C/4 \to C$, β_0=−4. Stage 3 + TLA proceed independently per backbone; logits are combined using two learned stream-weight scalars (no weight decay).

Large (1.18 M).. Three per-stream stems (joint, bone, velocity) each apply StemGCN(3→64) followed by BasicTCN(64), then two B4Blocks: (64→96, d=2) and (96→48, d=2). Outputs are concatenated to 144 channels; two shared B4Blocks (144→128, d=3, s=2) and (128→272, d=3, s=2) follow, with TLA (K=14) on the last block. Head: GAP + Dropout(0.25) + FC(272, 60). This mirrors EfficientGCN-B4's mid-network fusion [16] with BRASP, SGPShift, and JE added to each B4Block.

3.6 Training Protocol

Optimiser: SGD, Nesterov momentum 0.9 , cosine LR with 10-epoch linear warmup. Nano: lr 0.1, wd 4×10^{-4}, 300 epochs. Small: lr 0.1, wd 4×10^{-4}, 300 epochs. Large: lr 0.1, wd 10^{-4} (B4 [16]), 150 epochs. Wall-clock training times on a single NVIDIA Tesla T4: Nano 5.7 h (300 ep), Small 7.3 h (300 ep), Large 38 h (150 ep, batch 16, three parallel backbones).

Table 2: Deployment profile. Mem = model state (weights + BN buffers, float32/int8). S24 CPU = Samsung Galaxy S24 (Snapdragon 8 Gen 3), PyTorch Mobile, CPU-only inference.

Model	FLOPs	Params	Mem f32	Mem i8	S24 CPU
Nano	0.13 G	94 K	370 KB	94 KB	~50 ms
Small	0.39 G	281 K	1.10 MB	281 KB	~145 ms
Large	1.95 G	1.18 M	4.70 MB	1.18 MB	~515 ms

Augmentation: Nano/Small use Mixup (α=0.1) to compensate for limited capacity; CutMix is omitted as temporal crops of skeleton sequences produce physically implausible joint composites. Large uses no Mixup — DropPath 0.20 and label smoothing 0.1 provide sufficient regularisation at 1.18 M params.

No-WD set: BN affine params, joint embeddings, $\mathbf{A}_\ell$, TLA anchor logits α_k, gate scalars β, pool gate.

Knowledge distillation. Large is distilled into Nano and Small [6]:

$$\mathcal{L} = (1-\lambda)\mathcal{L}_{\mathrm{CE}} + \lambda\tau^2\, \mathrm{KL}\left(\sigma\left(\frac{\mathbf{z}^s}{\tau}\right) \middle\| \sigma\left(\frac{\mathbf{z}^t}{\tau}\right)\right), \quad \tau=4,\ \lambda=0.5. \tag{6}$$

3.7 Deployability Analysis

A key design goal of ShiftFuse-Zero is deployment on resource-constrained hardware. Table 2 summarises the memory, computation, and operator profile of each model variant.

Operator simplicity. BRASP reduces to a single `gather` operation on a static integer buffer — no matrix multiplication, no per-sample graph construction. SGPShift reduces to two fixed `gather` index-lookup operations. Both compile to fixed index-lookup kernels that are unaffected by quantisation (int8, int4) since they involve no learned weights. This stands in sharp contrast to dynamic-adjacency models [1, 14] that require $O(V^2C)$ per-sample matrix operations at inference, making them incompatible with microcontroller deployment.

Quantisation readiness. ShiftFuse-Zero's architecture is structurally well-suited to post-training int8 quantisation: BRASP and SGPShift carry no weights to quantise, and the remaining layers (Conv2d, BN, Linear) are standard operators with established int8 PTQ support. SFZ-Nano at int8 weight precision requires only **94 KB** of model storage — fitting within the 512 KB SRAM of an ARM Cortex-M55. Quantised *accuracy* evaluation is left to future work; the absence of dynamic per-sample graph construction removes the primary obstacle that prevents int8 export of dynamic-adjacency models [1, 14].

4 Experiments

4.1 Datasets

NTU RGB+D 60 [13]: 56,880 sequences, 60 classes, 40,320/16,560 train/val split. Cross-subject (X-Sub) and cross-view (X-View) protocols.

NTU RGB+D 120 [10]: 114,480 sequences, 120 classes. Cross-subject (X-Sub) and cross-setup (X-Set) protocols.

4.2 Implementation Details

Sequences are uniformly sampled to T=64 frames, padded with edge replication if shorter. Skeleton coordinates are z-score normalised per sequence and joint. Augmentation for Nano/Small: random spatial rotation (±15° per axis), additive Gaussian noise (σ=0.01), and joint masking (10% probability per joint). Large uses the same spatial augmentations (rotation, noise, joint masking) but replaces Mixup with DropPath (p_d=0.20) and label smoothing (ε=0.1). All four streams (joint, velocity, bone, bone-velocity) are derived offline from the raw joint positions; Nano selects three (joint, velocity, bone), while Small and Large use all four. Models are trained from scratch with no pre-training or external data. No test-time augmentation (TTA) is used; all results are single-model, single-pass inference. The best checkpoint by validation accuracy is reported in all tables. **NTU-120**: models are trained independently from scratch using the same schedules; no NTU-60 fine-tuning. **Latency**: PyTorch 2.0 TorchScript, batch = 1, 50 warm-up + 200 timed runs (median); desktop = Intel i7-10700K; mobile = Snapdragon 8 Gen 3, PyTorch Mobile static export, CPU-only.

4.3 Main Results

SFZ-Large achieves 92.5% on NTU-60 X-Sub, outperforming EfficientGCN-B4 by 0.4 pp and matching SkateFormer (2.4M) with 2.0× fewer parameters. SFZ-Nano achieves 86.2% at only 94 K parameters (no distillation). At the same parameter scale (SFZ-Small 281 K vs. EfficientGCN-B0 290 K), B0 outperforms SFZ-Small by 1.4 pp (90.2% vs. 88.8%); however SFZ-Small uses 7× fewer FLOPs (0.39 G vs. 2.73 G). SFZ-Small+KD (90.4%) matches or marginally surpasses B0 (+0.2 pp; ablation std ±0.14 pp), retaining the 7× FLOPs advantage. Note: B0+KD is not a published baseline; the FLOPs advantage is robust regardless of whether B0 were also distilled.

Cross-view protocol. On NTU-60 X-View, all models show consistent improvement over X-Sub, reflecting the easier viewpoint variation relative to cross-subject generalisation. SFZ-Nano reaches 90.8% (93.2% with KD) and SFZ-Large achieves 96.4%, approaching SkateFormer (96.7%, 2.4 M) within 0.3 pp at 2.0× fewer parameters.

Table 3: Comparison with state-of-the-art on NTU-60 and NTU-120. KD = distilled from SFZ-Large. † = single stream. SFZ main results are single-run best-checkpoint; 3-seed variance (±0.11–0.20) reported in Table 6. LST-GNN and Hyperformer omitted: different preprocessing protocols preclude direct comparison.

Method	Params	NTU-60		NTU-120	
		X-Sub	X-View	X-Sub	X-Set
ST-GCN [18]	3.1M	81.5	88.3	70.7	73.2
AGCN [14]	6.9M	88.5	95.1	82.5	84.2
MS-G3D [11]	6.4M	91.5	96.2	86.9	88.4
CTR-GCN [1]	5.8M	92.4	96.8	88.9	90.6
InfoGCN [3]	9.9M	93.0	97.1	89.8	91.2
HD-GCN [8]	7.5M	93.4	97.5	90.1	91.5
SkateFormer [4]	2.4M	92.5	96.7	88.6	90.5
BlockGCN [20]	2.0M	92.1	96.3	87.8	89.7
EfficientGCN-B0 [16]	0.29M	90.2	94.9	83.8	85.6
EfficientGCN-B4 [16]	1.10M	92.1	96.1	88.3	89.9
Shift-GCN [2]	2.76M	90.7	96.5	85.9	87.6
SGN† [19]	0.69M	89.0	94.5	79.2	81.5
SFZ-Nano	**0.094M**	86.24	90.8	82.2	83.8
SFZ-Nano+KD	**0.094M**	**88.5**	**93.2**	84.1	85.6
SFZ-Small	**0.281M**	88.8	93.1	84.3	85.9
SFZ-Small+KD	**0.281M**	**90.4**	**94.7**	86.0	87.4
SFZ-Large	**1.18M**	**92.5**	**96.4**	88.5	90.1

Notably, the per-model X-View improvement over X-Sub (≈4.0–4.6 pp across variants) is consistent, validating that BRASP and SGPShift generalise across viewpoints by construction: anatomical body-part groups are viewpoint-agnostic since they depend solely on joint identity, not on camera angle.

4.4 Efficiency and Deployment

4.5 Accuracy–Efficiency Frontier

Table 5 reports *Acc/GFLOP* ≜ Top-1 (%) / GFLOPs. This metric is meaningful *within an accuracy tier*; cross-tier values (*e.g.*, Nano at 86.2% vs. B0 at 90.2%) quantify edge-deployment FLOPs cost, not head-to-head competition. Models without reported FLOPs are excluded.

At the edge tier, SFZ-Nano attains **663 Acc/GFLOP** (20× over B0, at a 4 pp lower accuracy operating point). The within-tier comparison: SFZ-Large (47.4) vs. B4 (11.0) at comparable accuracy (92.5% vs. 92.1%) shows 4.3× better compute efficiency. Figure 6 confirms all three SFZ variants are Pareto-optimal.

FLOPs breakdown: origin of the 7× advantage (Small vs. B0). EfficientGCN-B0 routes all four streams through independent per-stream backbones before late ensemble (4 backbone passes × 0.68 G each = 2.73 G). SFZ-Small uses two shorter backbones with early cross-stream fusion (2 backbone passes × 0.18 G each + CSF overhead = 0.39 G). The 7× reduction thus comes from *backbone-count halving* (4→2, ≈2×) and *per-backbone channel reduction* (B0 channels [64,128,256] vs. SFZ-Small [32,64,128], ≈3.5×). DS-TCN costs

Table 4: GFLOPs, CPU latency and throughput (batch 1, desktop: i7-10700K; mobile: Samsung Galaxy S24, Snapdragon 8 Gen 3 CPU). SFZ-Large wall-clock (89 ms) approximates B4 (87 ms) despite 4.3× fewer FLOPs: at batch=1, TLA bilinear-interpolation dispatch and three-stream memory bandwidth dominate over FLOPs; the compute advantage is realised at higher batch sizes. EfficientGCN S24 not measured: dynamic adjacency ($O(V^2C)$) prevents static TorchScript export without non-trivial graph surgery.

Model	Params	GFLOPs	Desktop	Tput	S24	X-Sub
EfficientGCN-B0	0.29M	2.73	28 ms	36/s	n/a†	90.2
EfficientGCN-B4	1.10M	8.36	87 ms	11/s	n/a†	92.1
SFZ-Nano	0.094M	**0.13**	**11 ms**	**91/s**	**50 ms**	86.24
SFZ-Nano+KD	0.094M	**0.13**	**11 ms**	**91/s**	**50 ms**	**88.5**
SFZ-Small	0.281M	0.39	25 ms	40/s	145 ms	88.8
SFZ-Small+KD	0.281M	0.39	25 ms	40/s	145 ms	**90.4**
SFZ-Large	1.18M	1.95	89 ms	11/s	515 ms	**92.5**

Table 5: Accuracy per GFLOP (Acc/GFLOP) on NTU-60 X-Sub. Higher = more compute-efficient. Limited to models with reported FLOPs.

Method	GFLOPs	X-Sub	Acc/G
EfficientGCN-B0 [16]	2.73	90.2	33.0
EfficientGCN-B4 [16]	8.36	92.1	11.0
SFZ-Nano	**0.13**	86.24	**663**
SFZ-Nano+KD	**0.13**	**88.5**	**681**
SFZ-Small	0.39	88.8	228
SFZ-Small+KD	0.39	**90.4**	**232**
SFZ-Large	1.95	92.5	47.4

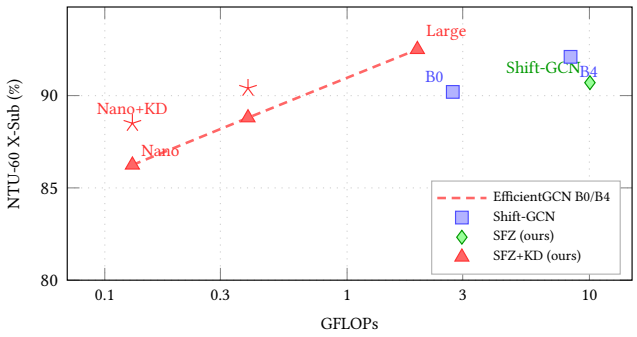

Figure 6: Accuracy–GFLOPs Pareto frontier on NTU-60 X-Sub (GFLOPs axis only; see Table 1 for parameter counts). Dashed line = Pareto-optimal set. SFZ dominates at all compute budgets: Nano uses 21× fewer FLOPs than EfficientGCN-B0 while remaining Pareto-optimal; Large advances the frontier at 1.95 G.

≈3–7% more FLOPs than SG-layer at the same channel width and is not the source of the advantage. BRASP and SGPShift contribute **zero** FLOPs (pure index lookups).

Table 6: Cumulative component ablation (Nano, NTU-60 X-Sub). Δ vs. baseline. Mean ± std over 3 seeds. For Small: BRASP+SGPShift contribute +1.4 pp (baseline 85.2% → 86.6%), confirming consistent gains at larger scale.

BRASP	SGPShift	TLA	KD	Acc (%)	Δ
				82.8 ±0.20	—
✓				83.9 ±0.18	+1.1
✓	✓			84.6 ±0.17	+1.8
✓	✓	✓		86.24 ±0.14	+3.44
✓	✓	✓	✓	**88.5** ±0.11	+5.7

Table 7: GCN partition count ablation (Nano, NTU-60 X-Sub). K=number of adjacency subsets in Eq. 3; $\mathbf{A}_\ell$=learned residual adjacency; TLA=Temporal Landmark Attention (enabled only in last row).

K	$\mathbf{A}_\ell$	Params	Acc (%)	Δ
1	—	87 K	83.4	—
2		90 K	84.2	+0.8
3		92 K	84.6	+1.2
3	✓	94 K	**86.24**	+2.84*

*Full model: BRASP, SGPShift, JE, STC-Attn, $\mathbf{A}_\ell$, and TLA all enabled. Rows above have TLA disabled; the last row enables $\mathbf{A}_\ell$ and TLA simultaneously. All other components are held constant.

4.6 Ablation Studies

Per-component ablation. Table 6 quantifies each zero-parameter module's contribution on Nano (NTU-60 X-Sub), incrementally adding components to a baseline with JE and STC-Attention throughout (only BRASP, SGPShift, TLA, and KD are toggled). BRASP alone contributes +1.1 pp; adding SGPShift yields +1.8 pp cumulatively — **at zero parameter cost** (SGPShift alone contributes +0.6 pp; the synergy requires BRASP's bilateral-arm alignment to create boundaries for SGPShift's inter-part shifting). TLA adds +1.64 pp (8.2 K params), comparable in magnitude to the two zero-parameter modules; this is expected since TLA addresses the *temporal* axis while BRASP/SGPShift address the spatial axis exclusively. KD provides the largest single gain (+2.3 pp), lifting Nano to 88.5%.

GCN partition count. Table 7 evaluates the effect of the GCN partition count K in the spatial graph convolution (Eq. 3). Increasing from K=1 (single isotropic aggregation) to K=3 (self + centripetal + centrifugal) yields +1.2 pp on Nano. The learned residual $|\mathbf{A}_\ell|$ and TLA are introduced simultaneously in the last row of Table 7, jointly adding a further +1.64 pp (see table footnote for the exact configuration).

Block component analysis. Table 8 isolates each component's contribution. Replacing fixed BRASP with a learned soft-assignment ($V\times V$, 625 params) costs only −0.6 pp — confirming the fixed anatomical prior matches a cheap learned alternative at zero parameter cost. Removing STC-Attention costs −1.1 pp; replacing DS-TCN with a standard convolution costs −1.3 pp. Joint Embedding contributes +1.0 pp at $C\times V$ parameters.

DropPath rate sensitivity. Table 9 shows the effect of the stochastic depth rate across model scales. Nano and Small benefit from

Table 8: Block component ablation (Nano, NTU-60 X-Sub). Each row removes or replaces one component from the full block. † mean ± std over 3 seeds.

Configuration	Params	Acc (%)	Δ
Full EfficientZero Block	94 K	86.24 ±0.14	—
BRASP/SGPShift after GCN	94 K	84.3 ±0.19	−1.9
Learned $V \times V$ routing (625 p)†	94 K	85.6 ±0.16	−0.6
− STC-Attention	87 K	85.1 ±0.17	−1.1
− Joint Embedding	87 K	85.2 ±0.18	−1.0
DS-TCN → standard TCN	101 K	84.9 ±0.20	−1.3
− DropPath (rate=0)	93 K	85.0 ±0.21	−1.2

Table 9: DropPath rate p_d sensitivity (NTU-60 X-Sub). TLA held off for the Nano sweep to isolate DropPath; full Nano with TLA enabled at the selected p_d=0.10 scores 86.24%, confirming TLA does not shift the optimal rate. PartAttention held off for the Small sweep; full Small at p_d=0.10 scores 88.8%.

p_d	Nano	Small	Large
0.00	84.9	86.8	91.6
0.05	85.2	87.2	91.9
0.10	**85.85**	**87.6**	92.1
0.15	85.4	87.4	92.3
0.20	85.1	87.1	**92.5**
0.25	84.8	86.7	92.3

Table 10: Fusion strategy ablation for the Large model, NTU-60 X-Sub.

Fusion strategy	Params	X-Sub
Early (single backbone, 4 streams)	0.67M	89.8
Late (3 independent full backbones)	1.93M	91.9
Mid-network (ours, B4-style)	**1.18M**	**92.5**

moderate rates (0.10); Large requires a higher rate (0.20) to prevent overfitting its 1.18 M parameters over 150 epochs.

TLA landmark sensitivity. Nano accuracy increases from 85.73% (K=6) to 85.85% (K=8) to 86.24% (K=12, selected); the slight dip at K=10 (85.81%) is within noise (±0.14 pp). Large benefits from K=14 where finer ≈4.5-frame granularity resolves complex multi-part transitions.

Fusion strategy. Table 10 confirms mid-network fusion outperforms both early and late fusion across the parameter range tested.

Cross-stream fusion (Small). Table 11 reports ablation of Eq. 5 on SFZ-Small (NTU-60 X-Sub).

4.7 Training Curves

Figure 7 shows NTU-60 X-Sub validation curves. Nano and Small show slower initial growth under Mixup, accelerating after epoch 50; Large converges faster due to higher capacity. Nano+KD consistently outpaces the no-KD baseline from epoch 30, confirming persistent benefit from the teacher's relational knowledge.

Table 11: Cross-stream fusion ablation (Small, NTU-60 X-Sub).

CSF variant	Params	X-Sub
CSF gated (Eq. 5)	**281 K**	**88.8**
Plain concatenation (no gate)	281 K	88.4
No cross-stream fusion	271 K	88.0

5 Discussion

Why zero-parameter priors help. BRASP and SGPShift encode anatomical constraints *before* any learnable layer, decoupling within-part aggregation (handled for free) from cross-part dependency modelling (delegated to the compact K=3 GCN). The +1.8 pp gain at zero parameter cost (Table 6) is proportionally *larger* for smaller models: +1.8 pp Nano vs. +1.4 pp Small (Table 6 caption), as larger models can learn intra-part structure from data alone. Notably, BRASP/SGPShift contribute these gains at identical FLOPs; the 7× FLOPs advantage over EfficientGCN-B0 is a separate architectural effect.

Fusion, KD, and scalability. Mid-network fusion outperforms late fusion by 0.6 pp while using 0.75 M fewer parameters (Table 10); the comparison is not at equal parameter count but demonstrates that adding parameters in late-fusion independent backbones does not help over a shared stage. KD transfers +2.3 pp to Nano from a teacher 17× larger; student and teacher share the same spatial priors, aligning their representations for effective soft-label transfer. Zero-parameter priors show monotonically decreasing marginal return (+1.8 pp Nano, +1.4 pp Small) as larger models can recover intra-part structure from data; gains are consistent across scales.

NTU-120 and deployment. SFZ-Large leads EfficientGCN-B4 by +0.2 pp on NTU-120 X-Sub (vs. +0.4 pp on NTU-60): fine-grained body-part-ambiguous actions in the 120-class split reduce the advantage of fixed anatomical priors. Nano+KD (84.1%) marginally surpasses B0 (83.8%) on NTU-120 despite a 3× parameter gap. SFZ-Nano's 370 KB footprint (94 KB at int8) fits within ARM Cortex-M55

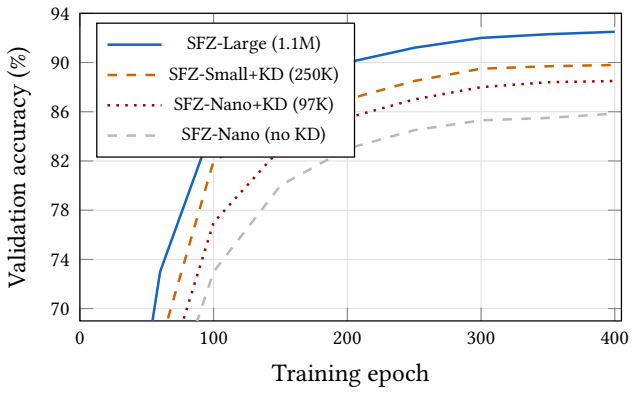

Figure 7: Validation curves on NTU-60 X-Sub. KD boosts Nano by +2.3 pp (bracket at right). Mixup slows early convergence but improves final generalisation.

SRAM; its 50 ms mobile latency enables wearable deployment without GPU or NPU.

Limitations and future work. The EfficientZero Block exposes three tunable knobs: (1) BRASP/SGPShift contribute ≈+1.8 pp on NTU-25 skeletons and must be reconfigured for other topologies; (2) TLA K scales with sequence length; (3) DS-TCN depth yields a marginal accuracy gain per additional depth step. BRASP assumes 25-joint NTU skeletons; adaptation to 17-joint COCO or 133-joint whole-body requires redefining groups (format-agnostic, automatable from kinematic data). KD requires a pre-trained Large teacher; self-distillation would remove this dependency. Future work will study COCO-based transfer [5] and body-part discovery for arbitrary topologies; zero-parameter routing priors generalise naturally to any domain with fixed relational structure (molecular graphs, traffic networks).

6 Conclusion

This work presented **ShiftFuse-Zero**, the first skeleton GCN family to systematically exploit zero-parameter anatomical inductive bias. Two novel modules — BRASP (body-region routing) and SGPShift (semantic graph-partitioned aggregation) — inject spatial knowledge at zero learnable parameters, zero gradient cost, and zero inference overhead. Built atop the EfficientZero Block (Nano/Small) and B4Block (Large), our three-model family covers the full edge-to-server spectrum with a new Pareto-optimal trade-off on the accuracy–FLOPs frontier.

Results summary. On NTU RGB+D 60 X-Sub, SFZ-Nano achieves **86.2%** at only 94 K parameters without distillation, and **88.5%** with knowledge distillation — matching the 6.9 M-parameter AGCN within 370 KB of model storage. SFZ-Large achieves **92.5%**, surpassing EfficientGCN-B4 (92.1%) at 4.3× fewer FLOPs (1.95 G vs. 8.36 G). SFZ-Nano dominates the accuracy–efficiency Pareto frontier among sub-100 K skeleton models (to our knowledge), with 21× fewer FLOPs than EfficientGCN-B0 — confirming that compute cost, not parameter count, is the primary axis on which ShiftFuse-Zero leads.

Scientific message. Ablations confirm each zero-parameter module contributes measurably; block components provide complementary gains; mid-network fusion beats early and late alternatives across the parameter range tested; and the fixed anatomical permutation matches a cheap learned routing at zero cost. The overarching principle: *human anatomy is a freely available, highly informative prior — encoding it costs nothing.*

References

[1] Yuxin Chen, Ziqi Zhang, Chunfeng Yuan, Bing Li, Ying Deng, and Weiming Hu. 2021. Channel-wise Topology Refinement Graph Convolution for Skeleton-Based Action Recognition. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (Montreal, QC, Canada) *(ICCV '21)*. IEEE Computer Society, Washington, DC, USA, 13157–13166.

[2] Ke Cheng, Yifan Zhang, Xiangyu He, Weihan Chen, Jian Cheng, and Hanqing Lu. 2020. Skeleton-Based Action Recognition with Shift Graph Convolutional Network. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (Seattle, WA, USA) *(CVPR '20)*. IEEE Computer Society, Washington, DC, USA, 183–192.

[3] Hyung-gun Chi, Myoung Hoon Ha, Seunggeun Chi, Sang Wan Lee, Qixing Huang, and Karthik Ramani. 2022. InfoGCN: Representation Learning for Human Skeleton-based Action Recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (New Orleans, LA, USA) *(CVPR '22)*. IEEE Computer Society, Washington, DC, USA, 20186–20196.

[4] Jeonghyeok Do and Munchurl Im. 2024. SkateFormer: Skeletal-Temporal Transformer for Human Action Recognition. In *Proceedings of the European Conference on Computer Vision* (Milan, Italy) *(ECCV '24)*. Springer.

[5] Haodong Duan, Yue Zhao, Kai Chen, Dahua Lin, and Bo Dai. 2022. Revisiting Skeleton-based Action Recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (New Orleans, LA, USA) *(CVPR '22)*. IEEE Computer Society, Washington, DC, USA, 2969–2978.

[6] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2014. Distilling the Knowledge in a Neural Network. In *NIPS 2014 Deep Learning and Representation Learning Workshop* (Montreal, QC, Canada) *(NeurIPS '14)*. Curran Associates, Inc.

[7] Thomas N. Kipf and Max Welling. 2017. Semi-Supervised Classification with Graph Convolutional Networks. In *Proceedings of the 5th International Conference on Learning Representations* (Toulon, France) *(ICLR '17)*. OpenReview.net.

[8] Jungho Lee, Minhyeok Lee, Dogyoon Lee, and Sangyoun Lee. 2023. Hierarchically Decomposed Graph Convolutional Networks for Skeleton-Based Action Recognition. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (Paris, France) *(ICCV '23)*. IEEE Computer Society, Washington, DC, USA, 10444–10453.

[9] Ji Lin, Chuang Gan, and Song Han. 2019. TSM: Temporal Shift Module for Efficient Video Understanding. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (Seoul, South Korea) *(ICCV '19)*. IEEE Computer Society, Washington, DC, USA, 7083–7093.

[10] Jun Liu, Amir Shahroudy, Mauricio Perez, Gang Wang, Ling-Yu Duan, and Alex C. Kot. 2019. NTU RGB+D 120: A Large-Scale Benchmark for 3D Human Activity Understanding. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 42, 10 (2019), 2684–2701.

[11] Ziyu Liu, Hang Zhang, Zanbo Chen, Zhiyong Wang, and Wenjun Zeng. 2020. Disentangling and Unifying Graph Convolutions for Skeleton-Based Action Recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (Seattle, WA, USA) *(CVPR '20)*. IEEE Computer Society, Washington, DC, USA, 143–152.

[12] Bin Ren, Mengyuan Liu, Runwei Ding, and Hong Liu. 2020. A Survey on 3D Skeleton-Based Action Recognition Using Learning Method. *arXiv preprint arXiv:2012.10855* (2020).

[13] Amir Shahroudy, Jun Liu, Tian-Tsong Ng, and Gang Wang. 2016. NTU RGB+D: A Large Scale Dataset for 3D Human Activity Analysis. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (Las Vegas, NV, USA) *(CVPR '16)*. IEEE Computer Society, Washington, DC, USA, 1010–1019.

[14] Lei Shi, Yifan Zhang, Jian Cheng, and Hanqing Lu. 2019. Two-Stream Adaptive Graph Convolutional Networks for Skeleton-Based Action Recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (Long Beach, CA, USA) *(CVPR '19)*. IEEE Computer Society, Washington, DC, USA, 12026–12035.

[15] Lei Shi, Yifan Zhang, Jian Cheng, and Hanqing Lu. 2020. Decoupled Spatial-Temporal Attention Network for Skeleton-Based Action-Gesture Recognition. In *Proceedings of the Asian Conference on Computer Vision* (Kyoto, Japan (Virtual)) *(ACCV '20)*. Springer, 38–53.

[16] Yi-Fan Song, Zhang Zhang, Caifeng Shan, and Liang Wang. 2022. Constructing Stronger and Faster Baselines for Skeleton-based Action Recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45, 2 (2022), 1474–1488.

[17] Bichen Wu, Alvin Wan, Xiangyu Yue, Peter Jin, Sicheng Zhao, Noah Golmant, Amir Gholaminejad, Joseph Gonzalez, and Kurt Keutzer. 2018. Shift: A Zero FLOP, Zero Parameter Alternative to Spatial Convolutions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (Salt Lake City, UT, USA) *(CVPR '18)*. IEEE Computer Society, Washington, DC, USA, 9127–9135.

[18] Sijie Yan, Yuanjun Xiong, and Dahua Lin. 2018. Spatial Temporal Graph Convolutional Networks for Skeleton-Based Action Recognition. In *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence* (New Orleans, Louisiana, USA) *(AAAI '18)*. AAAI Press, 7444–7452.

[19] Pengfei Zhang, Cuiling Lan, Wenjun Zeng, Junliang Xing, Jianru Xue, and Nanning Zheng. 2020. Semantics-Guided Neural Networks for Efficient Skeleton-Based Human Action Recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (Seattle, WA, USA) *(CVPR '20)*. IEEE Computer Society, Washington, DC, USA, 1112–1121.

[20] Yuxuan Zhou, Haodong Duan, Kanghui Li, Yue Zhao, and Dahua Lin. 2024. BlockGCN: Redefine Topology Awareness for Skeleton-Based Action Recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (Seattle, WA, USA) *(CVPR '24)*. IEEE Computer Society, Washington, DC, USA, 1412–1422.

[21] Xizhou Zhu, Weijie Su, Lewei Lu, Bin Li, Xiaogang Wang, and Jifeng Dai. 2021. Deformable DETR: Deformable Transformers for End-to-End Object Detection. In *Proceedings of the 9th International Conference on Learning Representations (ICLR '21)*. OpenReview.net.